

Implementing a MQTT broker provides 3 crucial business needs :

- **Insight :**
 - The ability to monitor and understand the current situation.
- **Hindsight :**
 - The ability to understand effects of an event after it has happened, although you did not understand at the time.
- **Foresight**
 - AI and Machine Learning algorithms provide the ability to predict what will happen in the future based on collected and stored information.



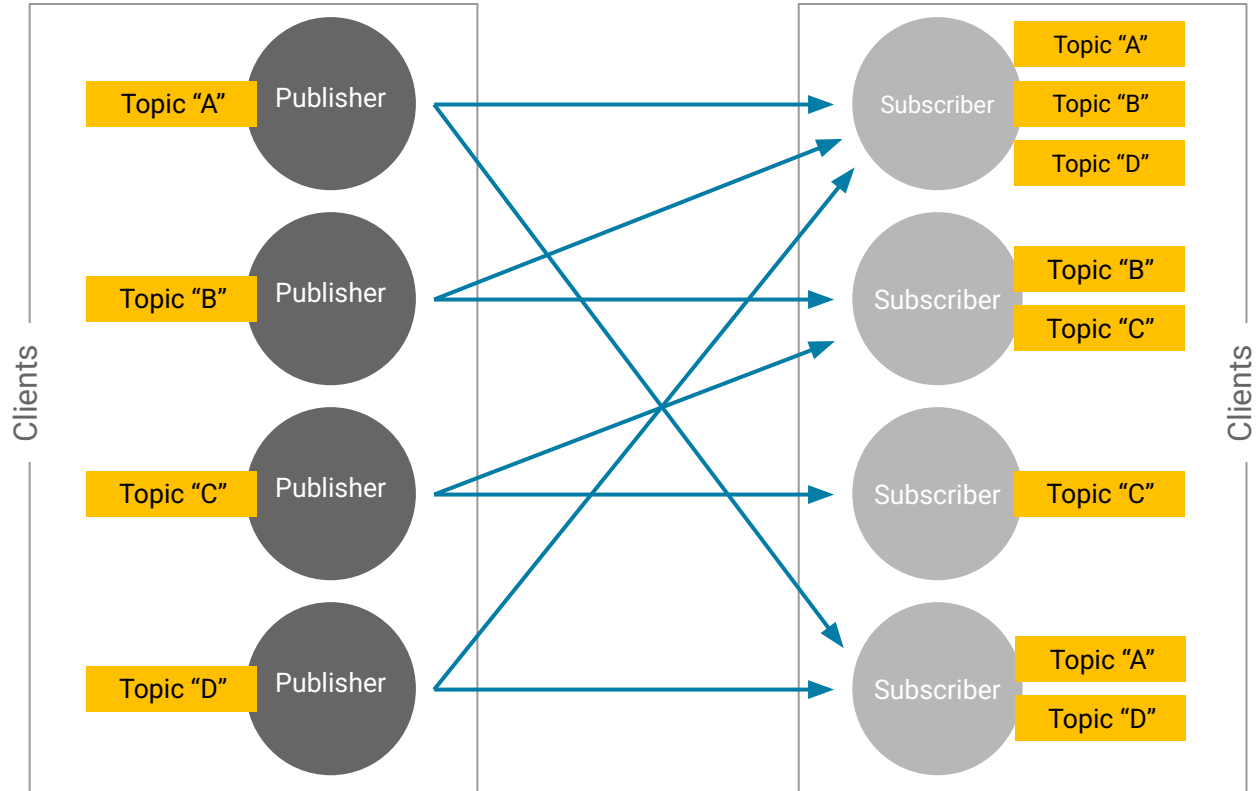
API's and OPU/UA: One on One Relationships

A lot of relation administration to administer

All publishers of data must “know” all Consumers; where they are and what data they produce.

Connection handling involves:

- Security
- Connection loss
- etc.



Insert the MQTT broker :

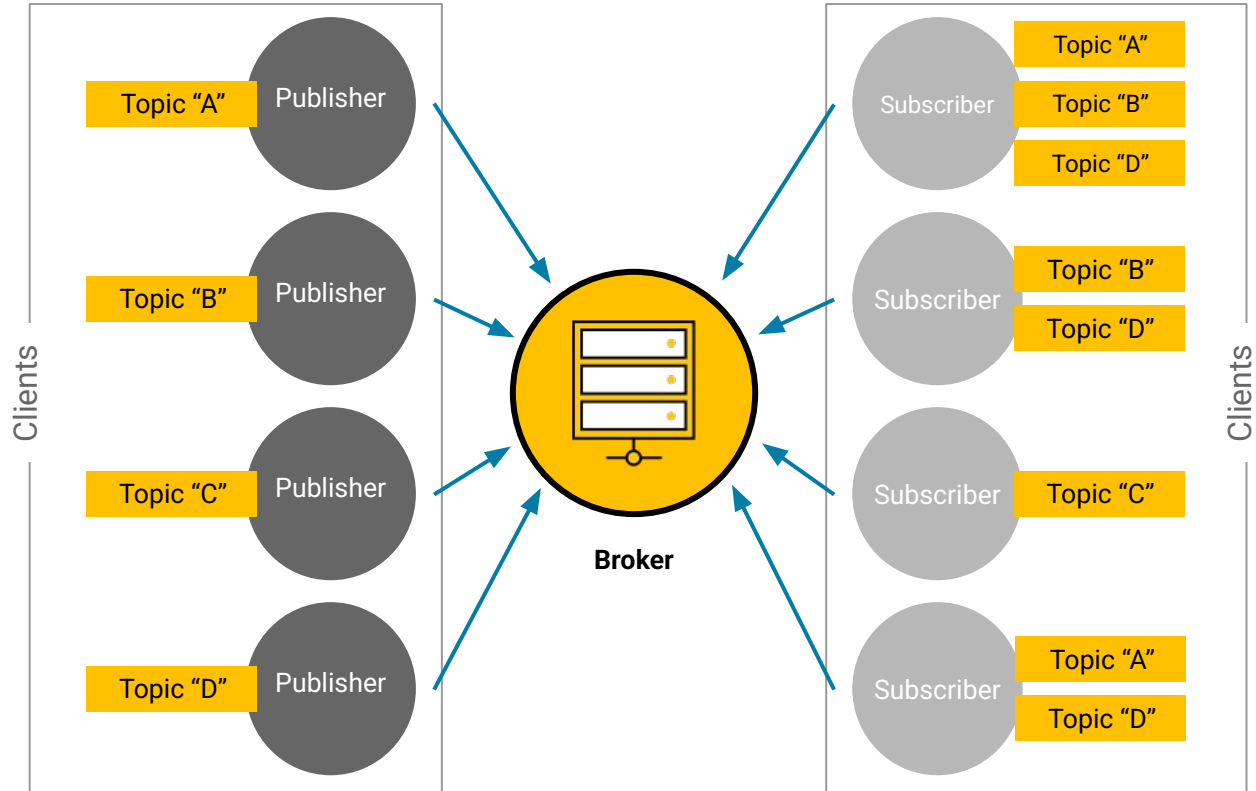
Less administration and more benefits

A client only needs to know:

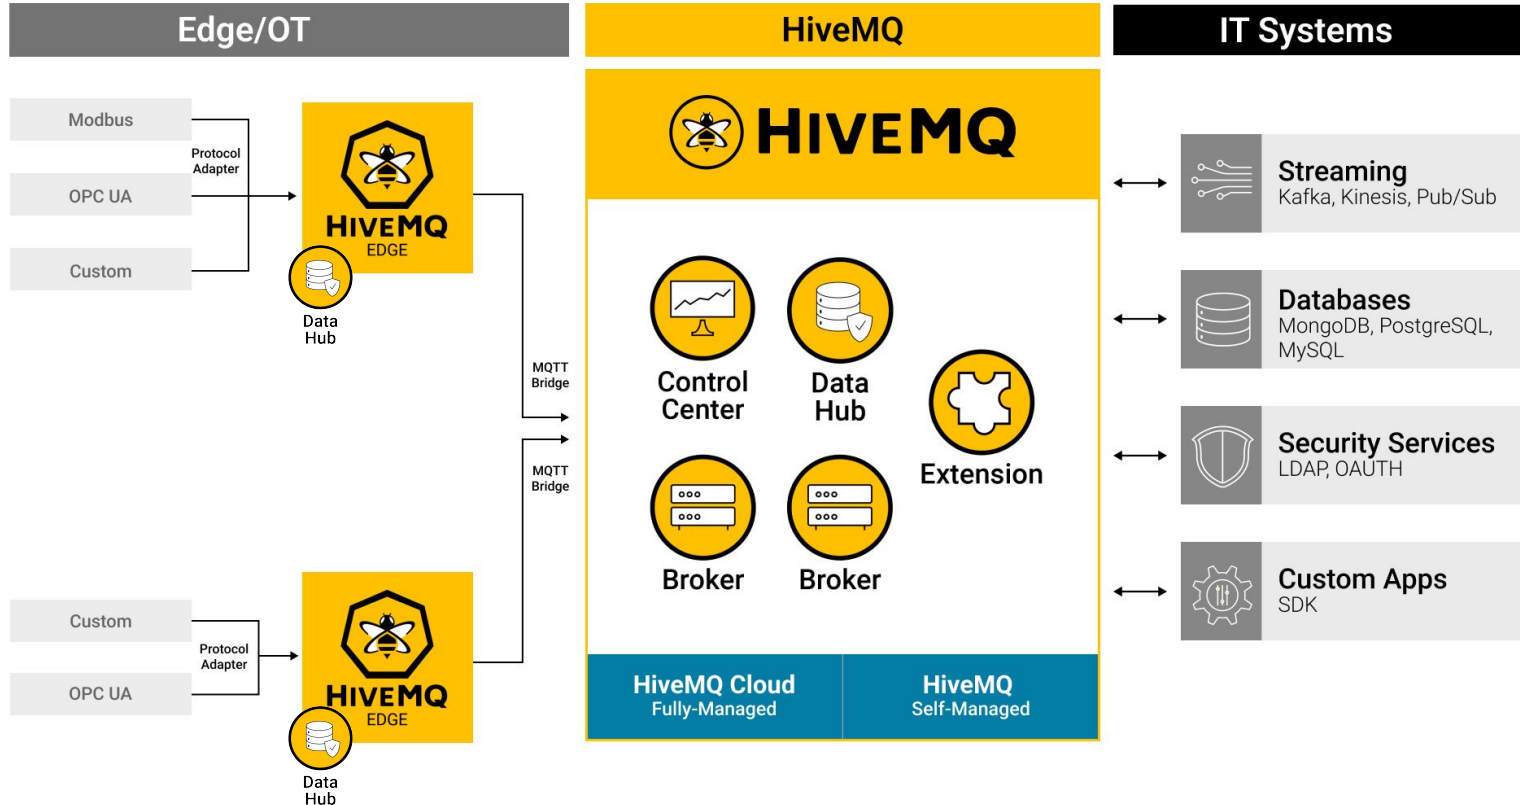
- A. The broker (IP)
- B. The topic to pub/sub to

Besides that the broker handles :

- Reconnects
- Disappearing Publishers (Last Will & Testament)
- Security
- and lots more



Edge & Enterprise MQTT Platform



“

HiveMQ with ESE, Swarm, Datahub, Kafka and Grafana

H01 - A dockerized hands-on session

Directory Structure



```
Hivemq /
├── bin (run.sh/bat)
├── conf
│   ├── examples
│   │   ├── configuration
│   │   └── logging
├── data
│   └── < STAY AWAY >
├── extensions
│   ├── hivemq-allow-all-extension - DISABLED
│   ├── hivemq-amazon-kinesis-extension
│   │   └── conf
│   ├── ..
│   ├── hivemq-bridge-extension
│   │   └── conf
│   └── hivemq-data-lake-extension
│       └── conf
├── license
├── log (hivemq.log)
├── tools
│   ├── hivemq-swarm
│   │   └── scenario
│   └── mqtt-cli
```



Extension Structure

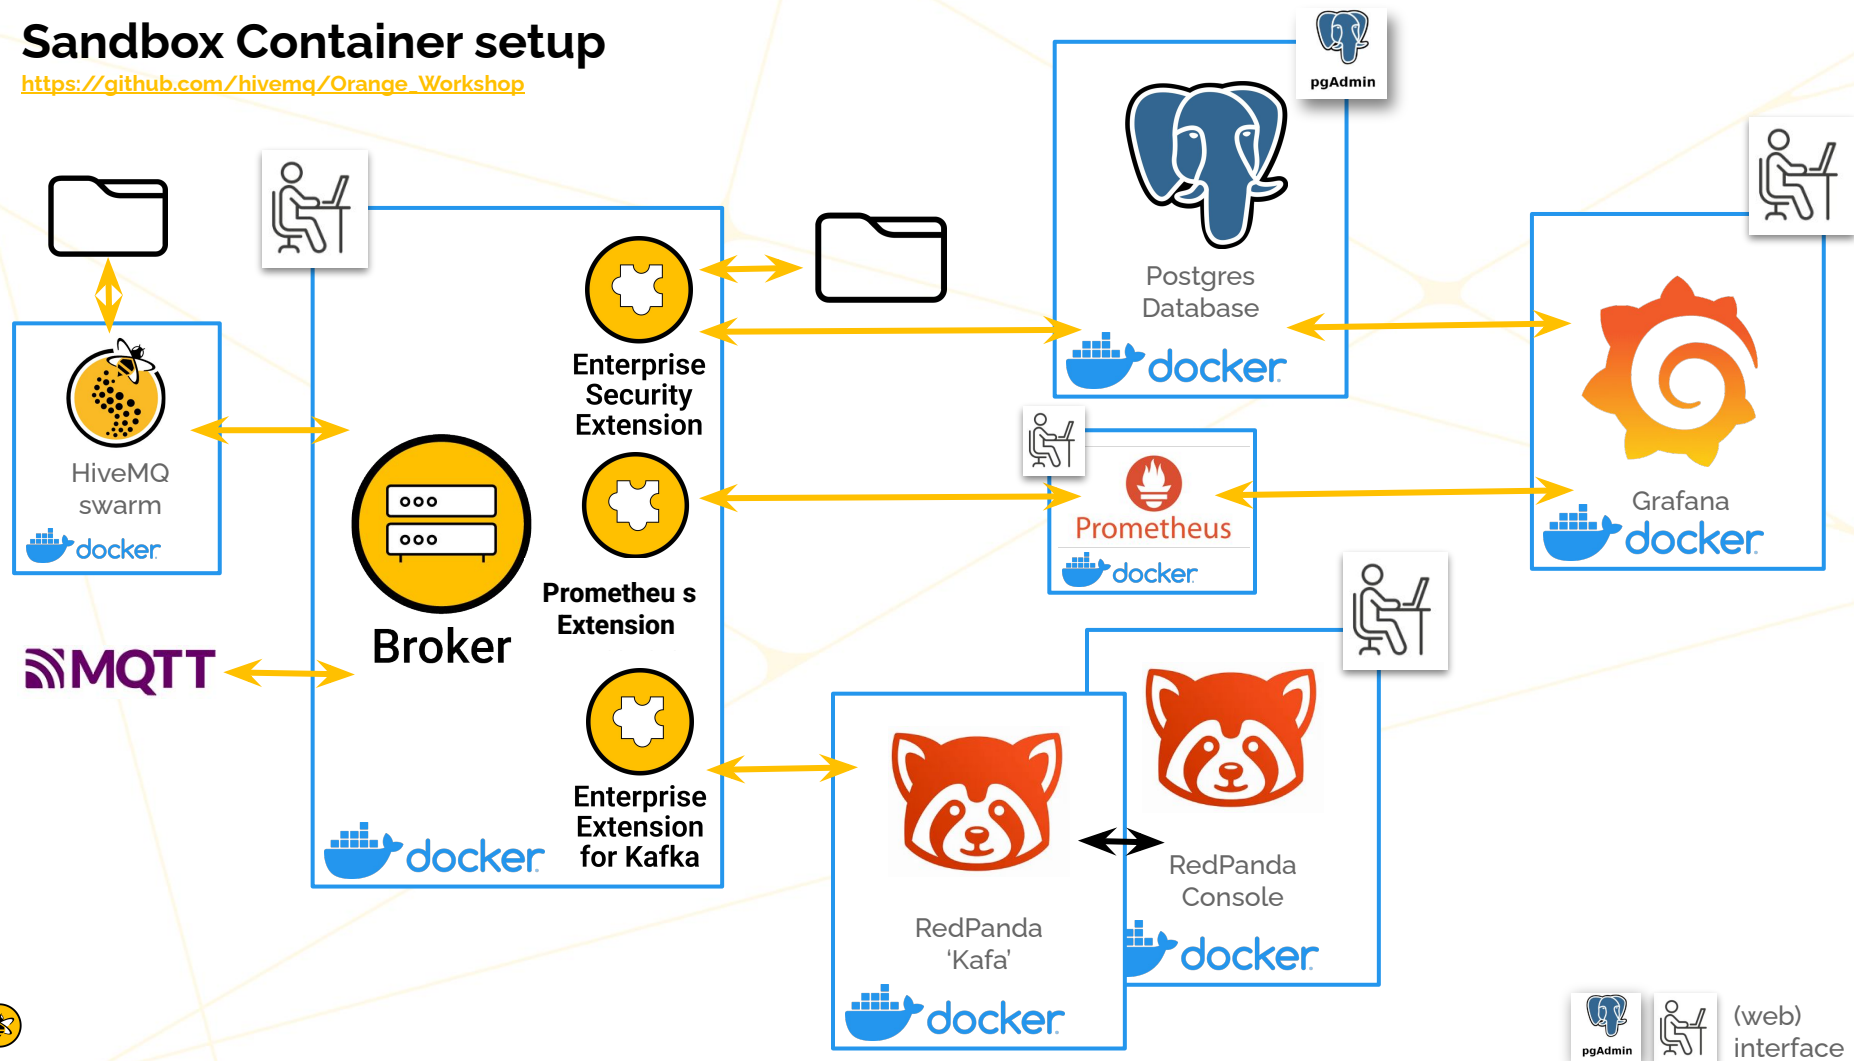


- Can be placed (or is already present) in the **extensions** directory.
- Connections and extensions are individually licensed. HiveMQ comes with a 25 connections and a 5 hr time-bombed ext set.
- All extensions are **DISABLED** with a flag.
- Default the **Allow-All** 'security' extension is enabled.
- All extensions are **HOT** reloadable. (Re)set the **DISABLED** flag.
- The **<ext.>/conf** directory is containing extension specific XML-based config (and sometimes support-) files.
- Check your **<HiveMQ>/log/hivemq.log** file.



Sandbox Container setup

https://github.com/hivemq/Orange_Workshop



HO-2 Enterprise Security Extension (ESE)

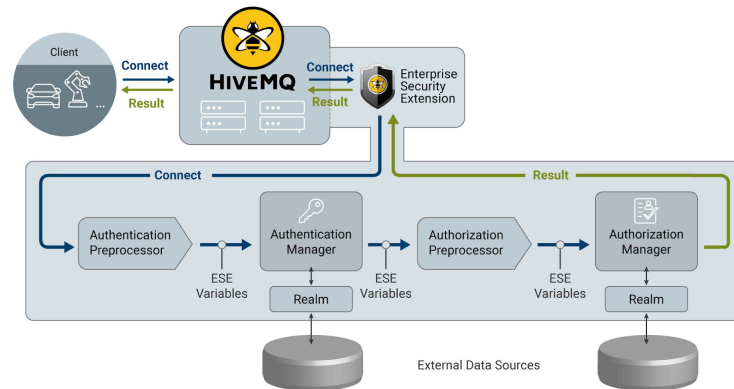
File based RBAC realm

```
<enterprise-security-extension
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">

  <realms>
    <file-realm>
      <name>file-realm</name>
      <enabled>true</enabled>
      <configuration>
        <file-path>conf/ese-file-realm.xml</file-path>
      </configuration>
    </file-realm>
  </realms>

  <pipelines>
    <listener-pipeline listener="ALL">
      <file-authentication-manager>
        <realm>file-realm</realm>
      </file-authentication-manager>
      <file-authorization-manager>
        <realm>file-realm</realm>
      </file-authorization-manager>
    </listener-pipeline>
  </pipelines>
</enterprise-security-extension>
```

XML



```
<user>
  <name>mqtt-user-1</name>
  <password>mqtt-password-1</password>
  <permissions>
    <permission>
      <topic>#</topic>
      <qos>ALL</qos>
      <activity>ALL</activity>
      <retain>ALL</retain>
      <shared-subscription>ALL</shared-subscription>
      <shared-group>group-1</shared-group>
    </permission>
  </permissions>
</user>
```

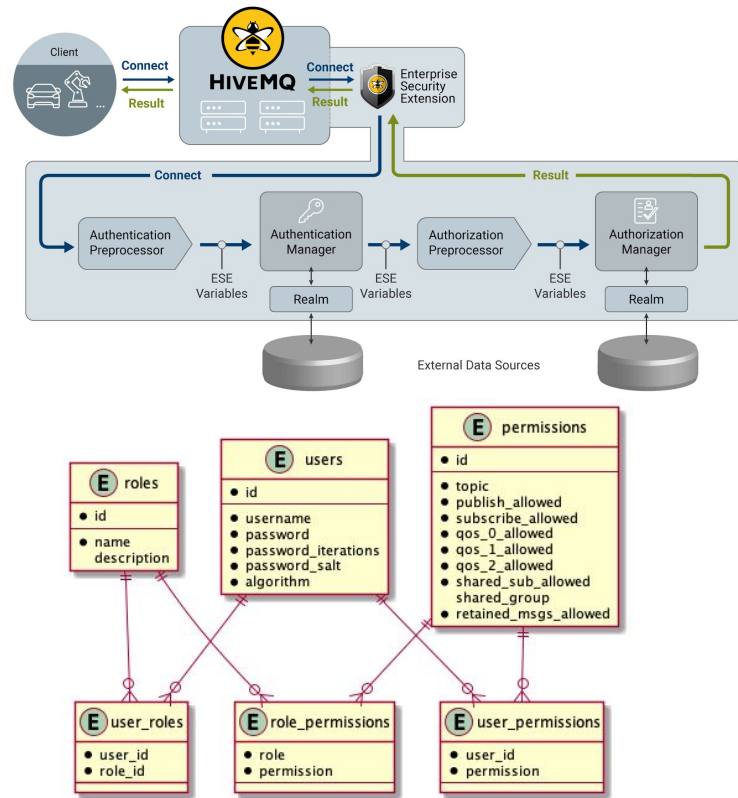
ese-file-realm.xml

H)2 - Enterprise Security Extension (ESE)

Postgres RBAC realm

```
<?xml version="1.0" encoding="UTF-8" ?>
<enterprise-security-extension
  xmlns:xsi="http://www.w3.org/2001/XMLSchema"
  version="1">
  <realms>
    <!-- a postgresql db-->
    <sql-realm>
      <name>postgres-backend</name>
      <enabled>true</enabled>
      <configuration>
        <db-type>POSTGRES</db-type>
        <db-name>hivemq_ESE</db-name>
        <db-host>postgres_db</db-host>
        <db-port>5432</db-port>
        <db-username>hivemq</db-username>
        <db-password>secret</db-password>
      </configuration>
    </sql-realm>
  </realms>
</enterprise-security-extension>
```

XML



H02 - Authorization of MQTT Clients

- ▶ Granular and flexible permission definitions using topic filters that can be reused for role and user permissions.
- ▶ **Filters** are built by the following expressions:
 - ▶ Preserved MQTT variables named **mqtt-clientid**, **mqtt-username** and **mqtt-password**
 - ▶ Use of general purpose variables e.g. **string-1**

Examples:

All topics with wildcard: **#**

Strict topic structure: **ohio/factory-01/machine-01/temperature**

Dynamic with MQTT Variables: **+/\${mqtt-username}/commands/#**

Dynamic with General Purpose Variables: **\${string-1}/\${mqtt-clientid}/commands/+**

- ▶ **Permission** that can be set:
 - ▶ **Method (Sub / Pub):** Enable publishing and subscription rights individually
 - ▶ **QoS:** Enable QoS levels individually
 - ▶ **Retained Messages:** Allow to send retained messages
 - ▶ **Shared Subscription:** Enable shared subscription rights and define allowed shared subscription group name



H03 - Postgres Extension

Uni-directional / Timescale, CrateDB etc.

```
<hivemq-postgresql-extension
xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance">

  <postgresqls>
    <postgresql>
      <id>my-postgresql-id</id>
      <host>postgres_db</host>
      <port>5432</port>
      <database>hivemq_ESE</database>
      <username>hivemq</username>
      <password>secret</password>
    </postgresql>
  </postgresqls>

  <mqtt-to-postgresql-routes>
    <mqtt-to-postgresql-route>
      <id>my-mqtt-to-postgresql-route-template</id>
      <postgresql-id>my-postgresql-id</postgresql-id>
      <mqtt-topic-filters>
        <mqtt-topic-filter>Temp-TS/#</mqtt-topic-filter>
      </mqtt-topic-filters>
      <processor>
        <statement-template>conf/statement.sql</statement-template>
      </processor>
    </mqtt-to-postgresql-route>
  </mqtt-to-postgresql-routes>
</hivemq-postgresql-extension>
```

XML

```
INSERT INTO tempdata (sensorid,isotime, unixtime,
temperature)
SELECT
  json_data->>'SensorID' AS sensorid,
  (json_data->>'isotime')::timestamp AS isotime,
  (json_data->>'unixtime')::numeric AS isotime,
  (json_data->>'temperature')::numeric AS temperature
/* casting from text to numic value ! */
FROM (
  VALUES
    ( ${mqtt-payload-utf8}::jsonb)
) AS input(json_data);
```

statement.sql

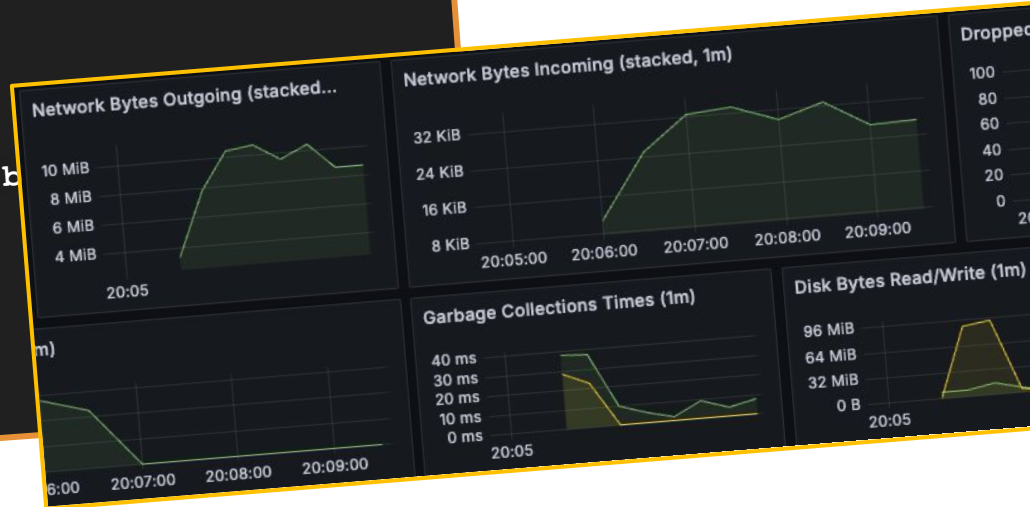


HO4 - Prometheus Extension


'User added'

```
# The ip where the servlet will be hosted  
ip=0.0.0.0  
# The port where the servlet will work on  
port=9399  
# The path for the servlet which gets called by  
# <ip>:<port> and <metric_path>  
# For example 127.0.0.1:9399/metrics  
metric_path=/metrics
```

prometheusConfiguration.properties



Prometheus
Extension



“

H06 - HiveMQ DataHub

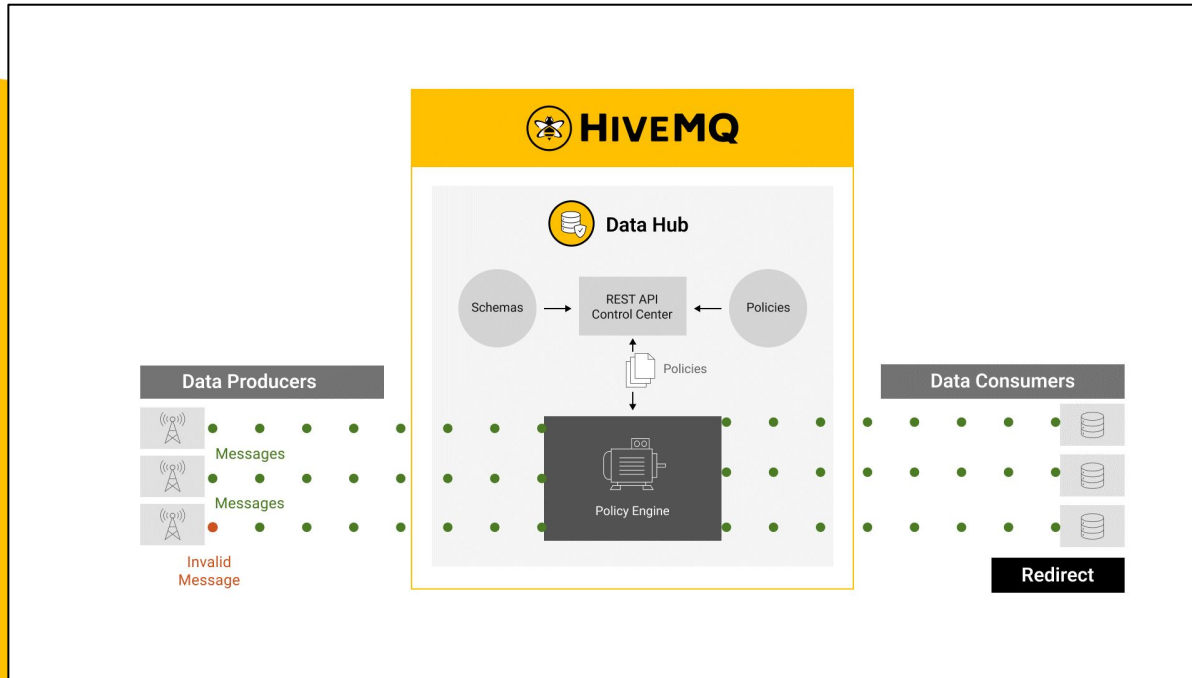


Data
Hub

Assure the quality of MQTT data:
Validate, contextualize and transform

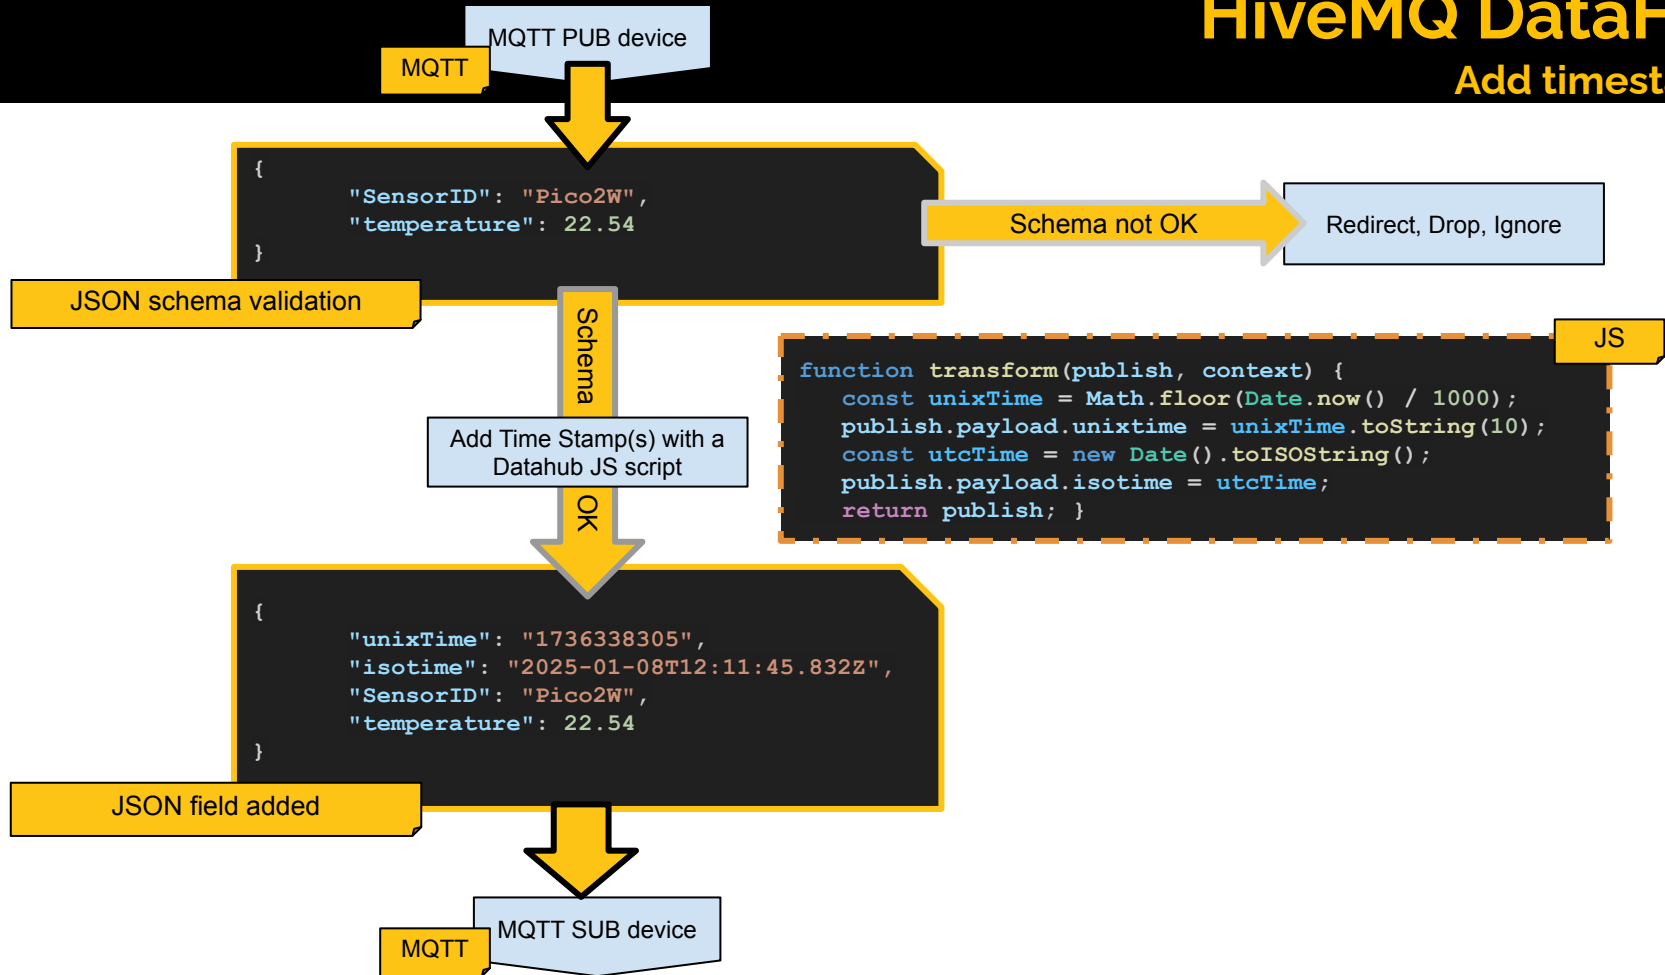
DataHub: How does it work?

Devices send data to the broker in a data format that is agreed upon by data producers and consumers and the broker enforces that format ensuring the consumer always gets the data in the expected format.



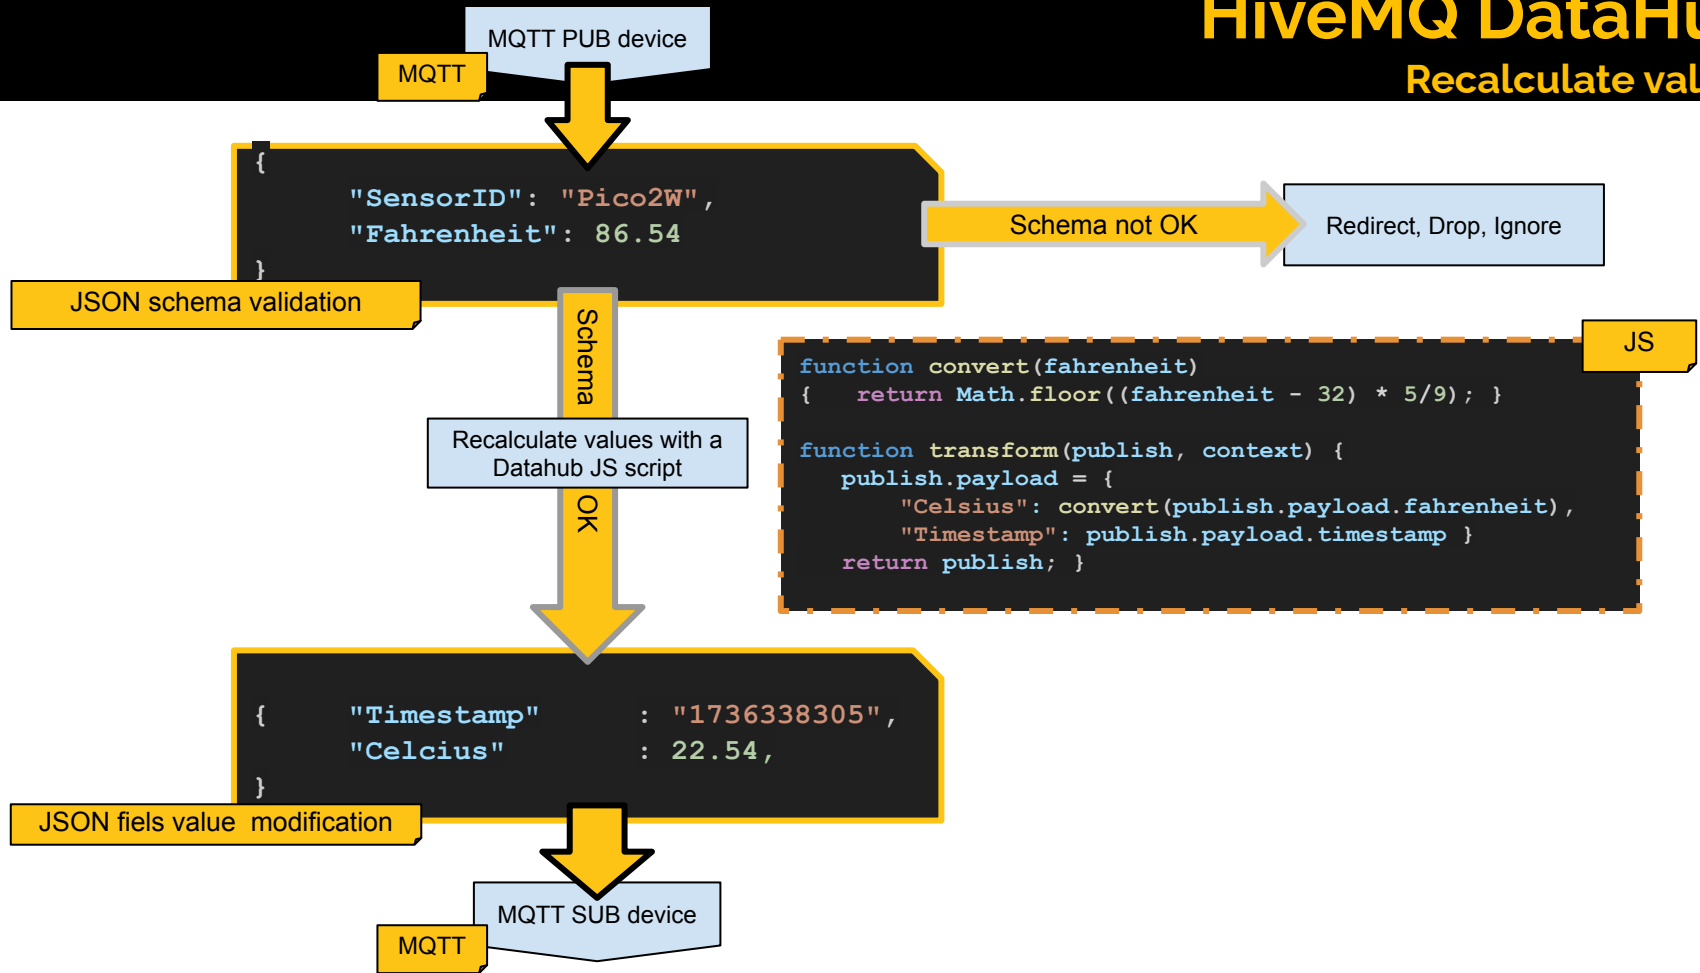
HiveMQ DataHub

Add timestamps



HiveMQ DataHub

Recalculate values



HO6b - Kafka Extension

Bi-directional / Azure Eventhub and Confluent !

```
<?xml version="1.0" encoding="UTF-8" ?>
<kafka-configuration
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:noNamespaceSchemaLocation="config.xsd">
  <kafka-clusters>
    <kafka-cluster>
      <id>cluster01</id>

      <bootstrap-servers>redpanda-0:9092</bootstrap-servers>
    </kafka-cluster>
  </kafka-clusters>
```

XML 1/2

```
<mqtt-to-kafka-mappings>
  <mqtt-to-kafka-mapping>
    <id>mapping 01</id>
    <cluster-id>cluster 01</cluster-id>
    <mqtt-topic-filters>
      <mqtt-topic-filter>to-kafka/#</mqtt-topic-filter>
    </mqtt-topic-filters>
    <kafka-topic>kafka-topic</kafka-topic>
  </mqtt-to-kafka-mapping>
</mqtt-to-kafka-mappings>

<kafka-to-mqtt-mappings>
  <kafka-to-mqtt-mapping>
    <id>mapping02</id>
    <cluster-id>cluster 01</cluster-id>
    <kafka-topics>
      <kafka-topic>test</kafka-topic>
    </kafka-topics>
    <kafka-topic-pattern>test-(.)*</kafka-topic-pattern>
  </kafka-to-mqtt-mapping>
</kafka-to-mqtt-mappings>
</kafka-configuration>
```

XML 2/2



HO6a -HiveMQ Swarm

Distributed, fully configurable MQTT load testing

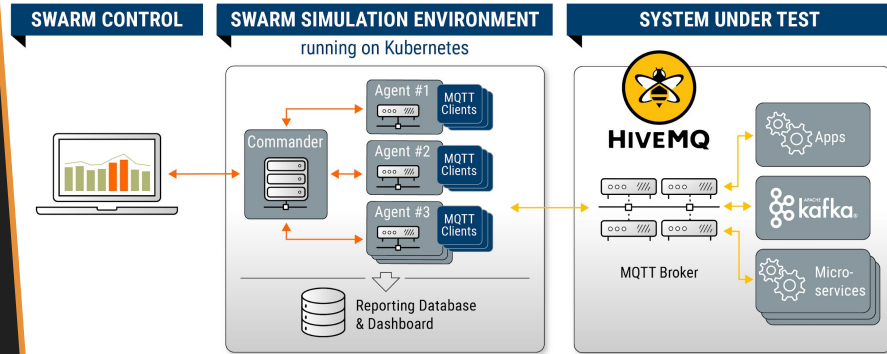
```
<brokers>
  <broker id="b1">
    <address>hivemq_broker</address>
    <port>1883</port>
  </broker>
</brokers>

<clientGroups>
  <clientGroup id="cg1">
    <clientIdPattern>A[0-9]{4}</clientIdPattern>
    <count>4</count>
  </clientGroup>
</clientGroups>

<topicGroups>
  <topicGroup id="tg1">
    <topicNamePattern>temp/subtopic-[0-9]</topicNamePattern>
    <count>3</count>
  </topicGroup>
</topicGroups>

<stages>
  <stage id="s1">
    <lifeCycle id="publish" clientGroup="cg1">
      <connect broker="b1"/>
      <publish topicGroup="tg1" count="100" message=' {
        "SensorID": "CMD-Line", "temperature": 22.54 } ' rate="1/5s"/>
    </lifeCycle>
  </stage>
</stages>
```

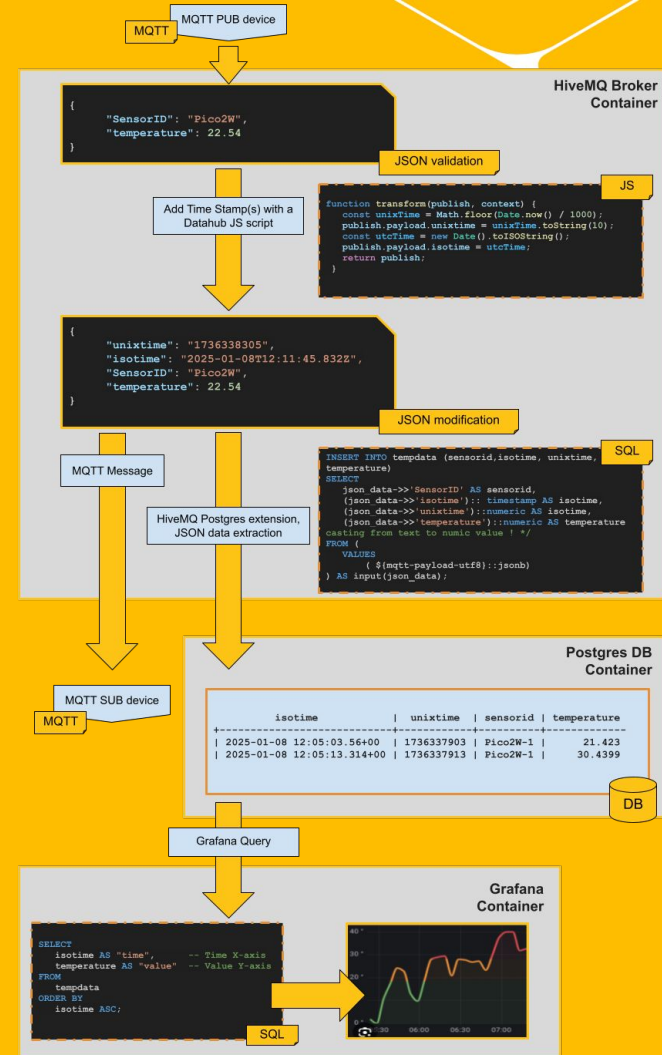
XML



HO6a - Usecase: End-2-End dataprocessing and visualisation

Starring ...

- HiveMQ Datahub
- HiveMQ Postgres ext.
- Postgres
- Grafana



Day 1 : Foundations of HiveMQ

Morning (8am to 12am)



Mod 1 -> Overview of HiveMQ: Understanding HiveMQ as an enterprise-ready platform.

Mod 2 -> Key Features and Architecture: Discuss the core features of HiveMQ, including scalability, security, and compliance with MQTT standards, extensions. Deep dive into the architecture of HiveMQ, focusing on its scalability and enterprise-grade security concepts.

Ho1 -> Hands-On: Basic configuration and setup of the Enterprise Broker.

Mod 3 -> Security Configuration: Configuring HiveMQ for enterprise-grade security, including authentication and authorization mechanisms.

Ho2 -> Hands-On: Add the Enterprise Security Extension to your deployment.

Afternoon (2pm to 7pm)



Mod 4 -> Enterprise Extensions: Detailed look at extensions for Streaming, Databases, and Data lakes.

Ho3 -> Hands-On: Add the Enterprise Database extension for PostgreSQL to your deployment.

Mod 5 -> Monitoring and Observability: Setting up monitoring for HiveMQ deployments.

Ho4 -> Hands-On: Add Prometheus extension to your deployment, collect the metrics and visualize with Grafana.

(Bonus) Mod 6 -> Performance Tuning: Techniques for optimizing HiveMQ performance in high-load scenarios and how to use Swarm to test it.

(Bonus) Ho5 -> Hands-On: Extend your deployment to a cluster and run a load test.



Day 2 : Data Hub and Edge

Morning (8am to 12pm)



Mod 7 -> HiveMQ Data Hub: Understanding Data Hub and its modules for data transformation and validation.

Ho6 -> Hands-On: Configure a policy and data transformation in your platform.

Mod 8 -> Introduction to HiveMQ Edge: Understanding HiveMQ Edge and the differences between the open source edition and the enterprise edition. Overview of pre-built protocol adapters (Beckhoff ADS, Ethernet IP, HTTP(s), OPC UA, ModBus, Siemens S7, File Input, Simulation, BACnet).

Ho7 -> Hands-On: Setup of HiveMQ Edge, Configuration of an OPC-UA protocol adapter with northbound and southbound mappings.

Afternoon (2pm to 7pm)



Mod 9 -> Our SDKs: Quick look at extensions and protocol adapter SDK and their capabilities.

Recap: Summary of the key points covered in the workshop.

Q&A Session: Open floor for questions and discussions.

Brainstorming : Share your challenges

